



Ratios of the mapping time  $T_{\text{mapping}}$  to the communication costs  $T_{\text{comm}}$  for HPC2 and Hadoop clusters for Dataset IV as a function of reciprocal cluster size  $1/N$ . Two HPC scenarios are shown as “HPC SLURM” and “HPC random”, which correspond to standard SLURM behaviour and a modified one, where nodes are being allocated from random racks. Linear fit done with the least-squares method.